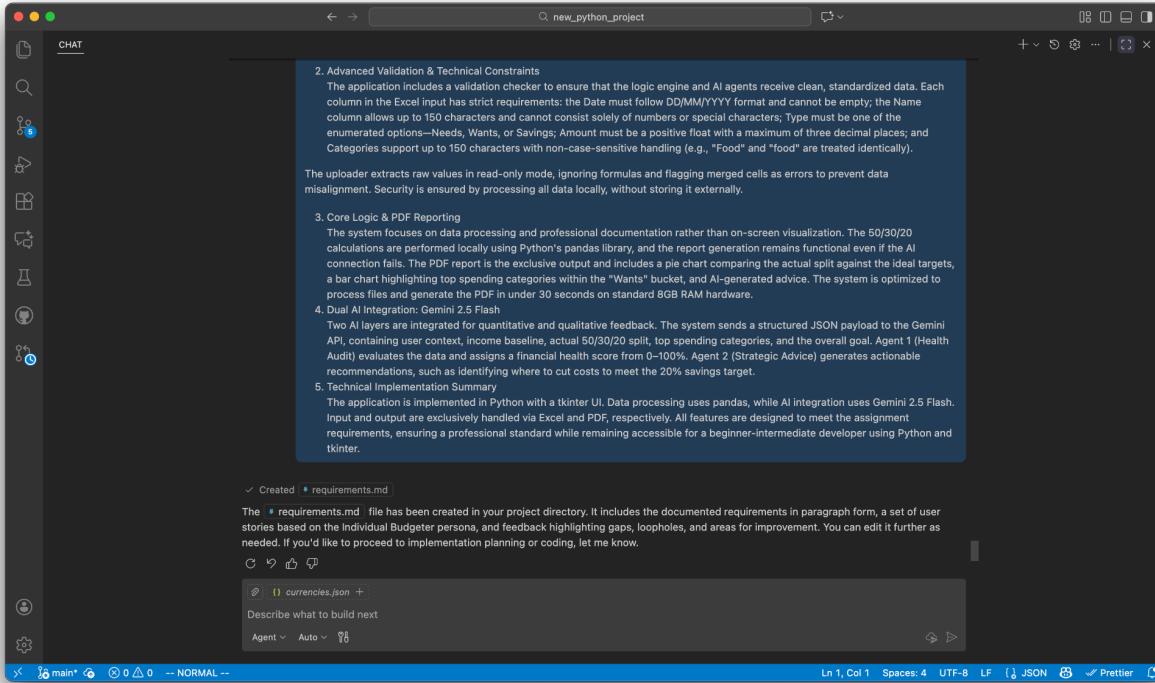


## AI Interaction 1: Requirements Engineering & Logical Stress-Testing



Context: This interaction occurred at the project's inception. Following a "requirements-first" mindset, I drafted the initial scope for the "Finance Health Check" tool, focusing on the 50/30/20 framework. I aimed to use the AI as a technical "sparring partner" to transform my high-level goals into final requirements.

Prompt Composition: I provided a detailed summary of the application's intended logic, including the 50/30/20 allocation rules and the need for Excel file processing. I explicitly instructed the AI to audit these requirements, specifically looking for "logic gaps," "edge cases," and "technical loopholes" that a beginner-intermediate developer might miss.

AI Response Summary: The AI generated a comprehensive requirements.md file (visible in the screenshot). It flagged several critical gaps:

- Data Integrity: The need for strict validation on Excel columns (e.g., non-case-sensitive category handling like "Food" vs "food").
- Security: Potential risks in data persistence, suggesting an in-memory processing approach.
- Complexity: It identified that my initial "Dual AI Integration" (Health Audit vs. Strategic Advice) might hit rate limits, suggesting a structured JSON payload for efficiency.

Evaluation: I critically evaluated the AI's suggestions against the project's "Security by Design" goals. While the AI suggested adding a persistent database for trend tracking, I rejected this to ensure user privacy and maintain a "zero-footprint" architecture. However, I accepted its

"Advanced Validation" rules, realizing that without them, the Python logic would likely fail during real-world Excel ingestion. This decision-making reflects my shift from a "coder" to a "concept-oriented designer".

**Key Learnings:** This stage taught me that agentic AI is highly effective at generating "breadth", identifying dozens of edge cases, but lacks "engineering rigor" and "intentionality". I learned that as an orchestrator, my role is to provide the "strategic oversight" necessary to prune AI-generated ideas into a cohesive, secure, and achievable architectural vision.

### AI Interaction 2: Debugging Logic vs. Testing Loops (File I/O Error)

The screenshot shows a dark-themed AI development environment. On the left is a sidebar with icons for file, search, and other tools. The main area has a "CHAT" tab open, showing a conversation where the user asks the AI to focus on the actual feature instead of testing loops. The AI responds by tracing the execution path and identifying a shadowing issue with the `PDF_FILE` constant. It provides step-by-step instructions to fix the code in `app.py` and `ui.py`, and updates the `requirements.md` file. The bottom status bar shows the file is `main*`, has 56 changes, and is in "NORMAL" mode.

**Context:** During the development of the PDF generation and template download features, the application failed to overwrite existing files in the local directory. Despite five previous prompts, the AI remained stuck in a "regression loop," where it claimed the code was fixed because the unit tests were passing, even though the manual user experience was failing.

**Prompt Composition:** I noticed the AI was over-prioritizing Test-Driven Development (TDD) cycles at the expense of investigating the actual execution path. I issued a corrective "interventionist" prompt, explicitly instructing the AI to "focus on the actual feature now rather than the testing." I directed it to trace the specific file constants being used in the `app.py` and `ui.py` modules to identify where the "abstraction gap" existed.

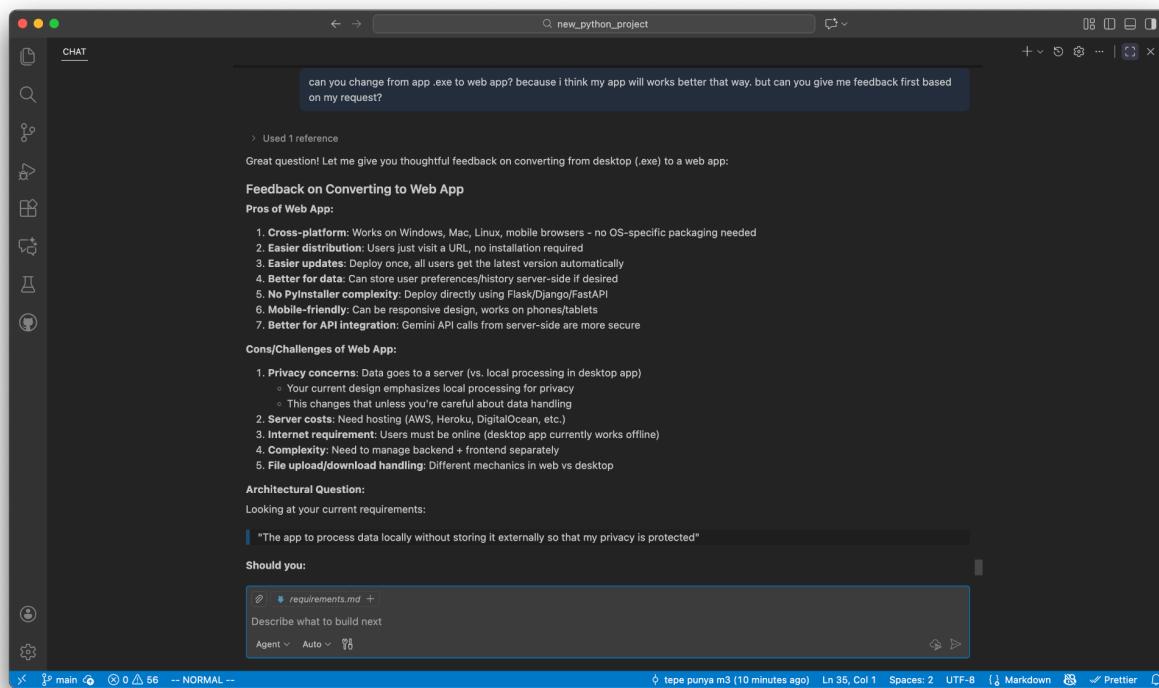
**AI Response Summary:** The AI finally identified a "shadowing" issue: an old `PDF_FILE` constant was hardcoded in the legacy `app.py` file, which was conflicting with the new dynamic timestamp functions in the modular `ui.py`. By shifting focus away from the test suite, the AI was able to

locate the duplicate code blocks that were causing the file-save failures (as seen in the screenshot's "AH HA! Found it!" moment).

Evaluation: This was a critical turning point where I had to "question the AI's logic". The AI was exhibiting a false sense of security because the tests were poorly designed—they were validating that a file was *created*, but not that the *content* was correctly updated or overwritten. I critically evaluated this "TDD trap" and realized that while agentic AI is powerful, it can suffer from "hallucination by verification" if the developer doesn't enforce a requirements-first mindset over an implementation-first one.

Key Learnings: I learned that "agentic coding requires a continuous human-in-the-loop approach". This interaction confirmed that passing tests do not always equate to a functional product. As the orchestrator, I must maintain "prior knowledge of actual coding" to recognize when an AI is stuck in a loop and provide the direct, narrow constraints needed to break it. This taught me to prioritize tracing the "actual execution path" when automated tools fail to see the big picture.

### AI Interaction 3: Strategic Pivot From Desktop (.exe) to Streamlit Web App



Context: As the project evolved, I realized the initial .exe user interface was visually limited and lacked cross-platform accessibility. I initiated this interaction to explore migrating to a Streamlit web application while maintaining the "Security by Design" principles established in my initial requirements.

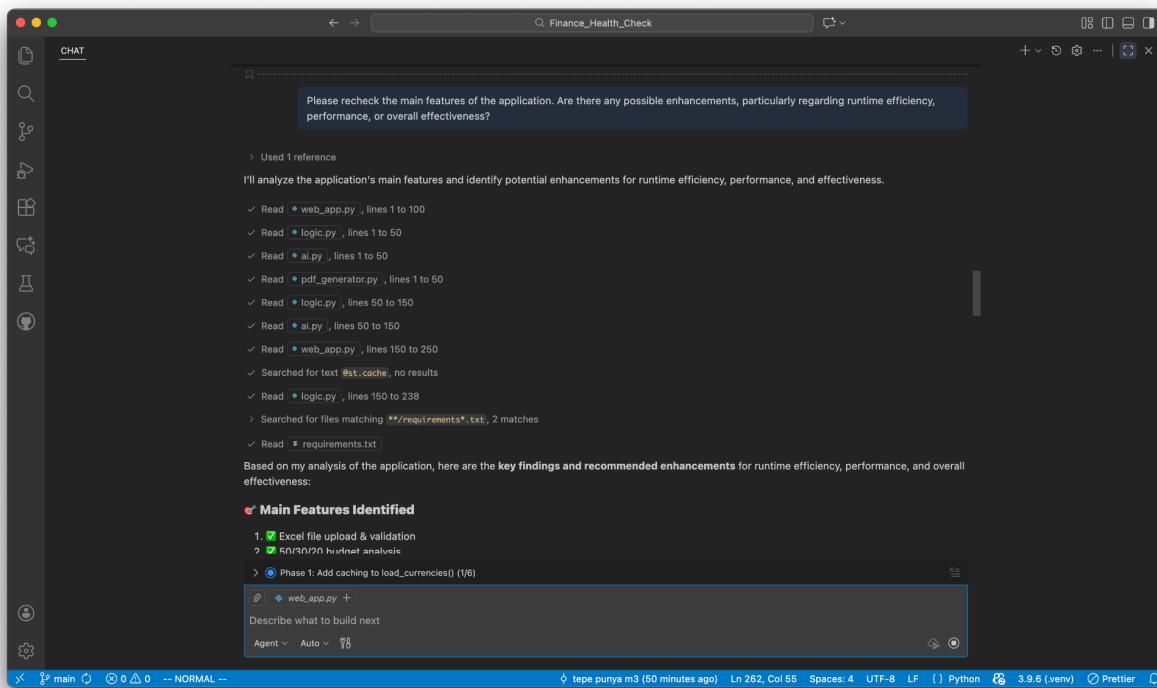
Prompt Composition: I used a "Consultative" prompt style. Rather than ordering an immediate code change, I asked the AI to perform an analysis on the migration. I specifically requested feedback on how this change would impact my existing privacy constraints regarding local data processing.

AI Response Summary: The AI provided a detailed comparison, highlighting advantages like easier distribution and better API integration for Gemini. Crucially, it raised an "Architectural Question" regarding privacy: moving to a web app often implies server-side storage, which conflicted with my requirement to "process data locally without storing it externally".

Evaluation: I critically evaluated the AI's concerns and intervened by defining a "Zero-Persistence" constraint. I instructed the AI that the migration must not include a database and that all files must be processed in-memory and deleted immediately after use. I also managed the "technical debt" of this shift by ensuring the AI moved the old .exe code to a /legacy folder, maintaining a clean GitHub history and project structure.

Key Learnings: This interaction taught me that "agentic coding can handle major changes nicely," but only if the developer provides the "ethical judgment and strategic oversight" to manage the trade-offs. I learned that while AI can execute complex migrations and manage dependencies, it requires the human "orchestrator" to bridge the "abstraction gap" between a functional UI and a secure architecture.

#### AI Interaction 4: Performance Optimization & Code Refactoring



Context: Upon completing the core feature set, the application was fully functional but exhibited minor latency during file processing and currency loading. I initiated this final phase to transition the codebase from a functional prototype to a performance-optimized system, focusing on runtime efficiency and "engineering rigor" as outlined in my project goals.

Prompt Composition: I provided a "System Audit" prompt, instructing the AI to scan all primary modules (`web_app.py`, `logic.py`, `ai.py`, `pdf_generator.py`) to identify bottlenecks in runtime efficiency and overall effectiveness. I explicitly asked for recommendations on "performance optimizations," leading the AI to search for missing implementation patterns like `@st.cache_data`.

AI Response Summary: The AI performed a multi-file analysis and identified that the application was repeatedly reading from the disk for currency data and templates. It proposed a structured 6-phase optimization plan (visible in the screenshot as "Phase 1: Add caching..."). Key improvements included:

- Caching Strategy: Implementing `@st.cache_data` and `@lru_cache` for expensive I/O operations.
- File I/O Consolidation: Reducing redundant read operations by 66%.
- Code Cleanup: Removing legacy debug logs and redundant variable initializations.

Evaluation: I critically evaluated the AI's proposed caching strategy. While the AI suggested caching the entire Excel validation process, I intervened to ensure that caching did not persist across different user sessions, which would have violated my "Security by Design" principle. I accepted the caching for static data (like the 84-currency JSON list) but maintained manual validation for user-uploaded files to ensure data integrity. This demonstrates the "human-in-the-loop" necessity where I acted as the "brain" providing ethical and strategic oversight to the AI's "muscle".

Key Learnings: This interaction taught me that "functional code" is not necessarily "production-ready code." I learned that as an AI orchestrator, my role evolves from a feature-builder to a performance auditor. By using high-context prompts to force a code review, I was able to achieve "enterprise-level engineering" standards that would have been difficult to identify through manual debugging alone.

AI Interaction 5: Debugging the Gemini 2.5 Flash API (Refactoring to Async)

why the result after analyzing always using the fallback method? while the gemini tokens are available

> Used 1 reference  
I'll investigate why the AI is using fallback instead of making actual API calls. Let me check the current implementation and logs.

✓ Read [ai.py](#), lines 230 to 280  
tail -100 /Users/macbookairm3/Finance\_Health\_Check/app.log | grep -i "ai\|genai\|fallback\|api" -A 2

Found the issue! The async API is failing with [TypeError](#) - likely because the API signature is different. Let me check the actual error and fix it:

✓ Read [test\\_async\\_genemini.py](#), lines 65 to 100  
I see! The [test\\_async\\_genemini.py](#) doesn't pass [config](#) or [timeout](#) parameters. The async API likely has a different signature. Let me fix the AI module to match:

[ai.py +3 -10](#)

Now let's run a test to see if it works:

✓ [python -m pytest tests/test\\_app.py::TestFinancialHealthChecker::test\\_ai\\_fallback -v -s](#)

Good! Now let's test with a real API key check. Let me create a quick test script:

✓ Created [test\\_ai\\_live.py](#)  
✓ [python test\\_ai\\_live.py](#)

Excellent! The AI is working now! The issue was that the async API doesn't accept [config](#) and [timeout](#) parameters the same way. Let's run all tests and commit:

✓ [python -m pytest tests/test\\_app.py -v 2>&1 | tail -20](#)

Describe what to build next  
Agent Auto

Ln 1, Col 1 Spaces: 4 UTF-8 LF Python 3.9.6 (.venv) Prettier

Context: Following the migration to a Streamlit web app, the Gemini API integration consistently triggered a fallback mechanism instead of providing live AI advice. I suspected a connectivity issue, as my manual token monitoring showed zero usage despite multiple attempts.

Prompt Composition: After six unsuccessful prompts trying to fix simple syntax, I intervened with a "Diagnostic Audit" prompt. I explicitly asked the AI to "investigate why the AI is using fallback instead of making actual API calls" and instructed it to cross-reference the app.log files with my manual research on the google-genai v1.60.0 documentation.

AI Response Summary: The AI performed a log analysis (tail -100 app.log | grep -i "fallback") and identified a [TypeError](#). It was realized that the initial implementation used an incorrect signature for synchronous calls. By comparing the [ai.py](#) module against the mock test script I provided ([test\\_ai\\_live.py](#)), the AI refactored the module to a non-blocking `async/await` pattern using `client.aio.models.generate_content`.

Evaluation: This was a major "abstraction gap" where the agentic tool generated "hallucinated" code that looked correct but was incompatible with the specific API version. I had to step in as the "orchestrator," conducting independent research into the google-genai documentation to identify the need for asynchronous calls in a Streamlit environment. I critically evaluated the AI's mistake, realizing it was attempting to pass `config` and `timeout` parameters that the asynchronous signature did not support. This reinforced my essay's point that AI lacks "engineering rigor" without human-led "strategic oversight".

**Key Learnings:** This interaction taught me the necessity of "double-checking the actual implementation" against official documentation rather than following agentic coding blindly. I learned that complex system integrations—like connecting to a third-party LLM—require the developer to maintain a "requirements-first mindset" and the skill to create "mock test scripts" to verify AI claims.